

UNDERSTANDING MULTI-TEACHER ON-POLICY DISTILLATION: FROM GRADIENTS TO CAPABILITIES

Anonymous authors

Paper under double-blind review

ABSTRACT

Multi-teacher on-policy distillation transfers domain expertise into a shared student. We study how loss averaging, optimizer history, and vocabulary supervision shape this transfer, using Qwen3-1.7B and four domain specialists, with additional SmoLLM3-3B measurements. Token and response averaging produce different gradients whose update directions become more similar after applying saved Adam moments; final task scores are close under top-64 supervision. Training changes about 97% of FP32 parameters, with only 7–9% remaining visibly changed after BF16 rounding. Adam’s accumulated moments also account for much of the alignment between teacher-specific updates: subtracting the zero-gradient step reduces their mean cosine from above 0.83 to below 0.01 on domain-specific inputs. Expanding supervision from a sampled token to the top-64 intersection improves final mathematics accuracy and lowers science accuracy, with relative gains changing during training. Top-64 closely approximates the full-vocabulary gradient direction in the high-coverage Qwen measurements; SmoLLM measurements with lower coverage retain appreciable gradient error. Together, these findings connect local update measurements to task-dependent learning trajectories and identify optimizer history and numerical precision as central to interpreting distillation updates.

1 INTRODUCTION

Multi-teacher on-policy distillation (MOPD) transfers specialized capabilities into a single model (Xiaomi LLM-Core Team, 2026; NVIDIA, 2026; Kimi Team et al., 2026; DeepSeek-AI et al., 2026). The student generates responses, and a teacher assigned to each prompt domain supplies token-level supervision. Improving several capabilities requires deciding how much each domain and response contributes to the loss, and how many next-token alternatives the teacher supervises.

Recent work establishes both the potential of MOPD and its sensitivity to these choices. Ma et al. (2026) study capability integration with policy-gradient and top- k objectives. Open-MOPD addresses uneven transfer through domain token balancing, performance-gap weighting, and refreshed rewards (Gao et al., 2026). The Nemotron 3 Ultra report finds no advantage for top- k or full-vocabulary matching over sampled-token supervision in preliminary experiments (NVIDIA, 2026, Section 3.3.5). Other studies report parameter changes concentrated in a small set of stored weights (Yu et al., 2026; Shen et al., 2026). How these observations connect loss construction, optimizer updates, and task performance remains open.

Our study follows the path from loss construction to gradients, parameter updates, and task performance. The main experiments train Qwen3-1.7B with teachers for mathematics, code, instruction following, and science. We vary loss averaging and vocabulary supervision during training, and compare gradients and optimizer steps at fixed student parameters and inputs. Three findings emerge:

- Adam brings different normalization gradients closer in update direction (Section 3).** Weighting responses by length or equally produces gradients with mean cosine 0.68; applying the same saved Adam state raises the update cosine to 0.96. Under top-64 supervision, the averaging rules reach similar final mean scores and follow different task-level learning curves.

2. **Precision and optimizer history strongly shape measured update sparsity and alignment (Section 4).** Training changes most FP32 weights, while BF16 rounding leaves a small fraction visibly changed. Teacher-specific steps are strongly aligned through shared Adam moments; subtracting the zero-gradient step leaves near-zero mean alignment on domain-specific inputs.
3. **Broader vocabulary supervision produces task-dependent gains (Section 5).** Top-64 supervision improves final mathematics accuracy and lowers science accuracy relative to sampled-token supervision. Its mean advantage narrows as training proceeds. At fixed student and optimizer state, the losses produce similar update concentration.

Additional SmoLLM3-3B measurements show that top-64 retains high average student probability while omitting substantial mass on some inputs. Enlarging the candidate set improves gradient approximation, with residual error relative to full-vocabulary supervision. Section 6 connects these observations to checkpoint selection: individual tasks can regress during training and recover by the endpoint, and development-set rankings can change on separate test questions. These results motivate tracking per-task learning curves alongside gradients and parameter changes.

2 PRELIMINARIES

Multi-teacher on-policy distillation (MOPD). For domain d , the student generates $y \sim p_\theta(\cdot | x)$ for $x \sim \mathcal{D}_d$, and a frozen teacher q_d supplies probabilities at each prefix $h_t = (x, y_{<t})$ (Gao et al., 2026). For a rollout batch $\mathcal{B}$, response i has T_i valid tokens and mean loss $\ell_i = T_i^{-1} \sum_t \ell_{i,t}$. The objective is

$$\mathcal{L}(\theta) = \sum_{i \in \mathcal{B}} w_i \ell_i(\theta), \quad w_i \geq 0, \quad \sum_{i \in \mathcal{B}} w_i = 1. \quad (1)$$

The token loss determines the vocabulary supervision; the response weights w_i determine normalization across tokens, responses, and domains. Each update uses fresh student responses, held fixed with their masks and weights during gradient computation.

Sampled-token, top- k , and full-vocabulary supervision. Write $p_\theta(v) = p_\theta(v | h_t)$ and $q_d(v) = q_d(v | h_t)$. Full-vocabulary supervision minimizes reverse KL over $\mathcal{V}$; the policy-gradient (PG) loss scores the sampled token y_t :

$$\ell_{\text{full}} = \sum_{v \in \mathcal{V}} p_\theta(v) \log \frac{p_\theta(v)}{q_d(v)}, \quad \ell_{\text{PG}}(y_t) = -\text{sg}[\log q_d(y_t) - \log p_\theta(y_t)] \log p_\theta(y_t) \quad (2)$$

Here sg denotes stop-gradient. The expected PG gradient equals the full reverse-KL gradient: differentiating Equation 2 gives $\mathbb{E}_{a \sim p_\theta} [\log(p_\theta(a)/q_d(a)) \nabla_\theta \log p_\theta(a)]$, since $\sum_v \nabla_\theta p_\theta(v) = 0$.

Top- k supervision uses $I = S \cap T$, where S, T are the student and teacher top- k sets, computed at rollout and held fixed. Its loss is

$$\ell_{\text{top-}k} = - \sum_{v \in I} \text{sg} \left[\frac{p_\theta(v)}{\sum_{u \in S} p_\theta(u)} (\log q_d(v) - \log p_\theta(v)) \right] \log p_\theta(v). \quad (3)$$

Probabilities use the full-vocabulary softmax; detached weights are normalized over S , and an empty intersection gives zero loss. I16 and I64 denote $k = 16$ and 64. We call the number of supervised vocabulary tokens the *supervision density*: one for PG, $|I| \leq k$ for intersection supervision, and $|\mathcal{V}|$ for full supervision.

Domain weighting and loss normalization. Let $\mathcal{B}_d$ contain responses from domain d , and let $\lambda_d \geq 0$ with $\sum_d \lambda_d = 1$ be the target domain loss weights. For nonempty responses and domains, the three rules assign response $i \in \mathcal{B}_d$ weights

$$w_i^{\text{GT}} = \frac{T_i}{\sum_{j \in \mathcal{B}} T_j}, \quad w_i^{\text{DT}} = \frac{\lambda_d T_i}{\sum_{j \in \mathcal{B}_d} T_j}, \quad w_i^{\text{DR}} = \frac{\lambda_d}{|\mathcal{B}_d|}. \quad (4)$$

Global-token (GT) averaging weights domains by their response-token shares. Domain-token (DT) and domain-response (DR) averaging fix each domain’s total weight to λ_d : DT weights responses by length, while DR weights their mean losses equally. Domain loss weights are distinct from prompt sampling proportions.

Table 1: **Task gains and method rankings vary across domains and evaluation sets.** Joint-training scores (%) after 500 updates. The four development tasks are MATH-500, LiveCodeBench v6, IFBench, and GPQA Diamond; both means weight domains equally. Response and domain-token averaging assign equal domain loss weights and average responses equally or by length; global-token averaging weights all tokens equally. Test questions are selected separately before inference.

Supervision	Averaging	Math	Code	IF	Science	Dev. mean	Test mean
Initial student	—	68.80	16.41	18.33	32.20	33.94	38.60
<i>Adam</i>							
Sampled token	Response	73.20	18.27	27.00	35.97	38.61	43.55
Sampled token	Domain token	73.20	19.24	27.00	34.96	38.60	43.51
Sampled token	Global token	76.20	17.92	27.33	32.94	38.60	44.51
Top-16	Response	75.40	17.75	27.00	35.46	38.90	44.56
Top-64	Response	75.80	18.30	27.33	34.71	39.04	44.46
Top-64	Domain token	74.00	18.86	28.67	35.22	39.19	43.68
Top-64	Global token	74.20	19.80	28.67	33.41	39.02	44.97
<i>SGD</i>							
Sampled token	Response	76.40	19.53	28.33	35.59	39.96	43.48
Sampled token	Domain token	75.20	20.31	28.33	33.90	39.44	43.90
Sampled token	Global token	75.00	19.26	25.33	34.71	38.58	42.97

Update similarity. Let u_d be the parameter change from one step on domain d ’s mean loss. Each compared step starts from identical student weights and optimizer state, including Adam’s moments. For nonzero updates, their cosine is $C_{de} = (u_d^\top u_e) / (\|u_d\| \|u_e\|)$.

3 LOSS NORMALIZATION AND TASK PERFORMANCE

Shared experimental setting. We train Qwen3-1.7B with four RL-specialized teachers derived from the same base model, covering mathematics, code, instruction following, and science. Joint training uses 500 updates with 16 fresh responses per domain per update; mathematics-only training uses the same total batch size. Comparisons within each optimizer share initialization, teachers, and prompt schedules. AdamW uses learning rate 2.5×10^{-7} , moments (0.9, 0.98), and norm-1 clipping. Responses are sampled at temperature 1 with a 4,096-token cap; thinking is disabled throughout.

Evaluation and uncertainty. Development scores use MATH-500, LiveCodeBench v6, IFBench, and GPQA Diamond, with 500/1,055/300/198 questions. The first three use greedy decoding; GPQA averages correctness over four sampled responses per question. Evaluation uses a 32,768-token context, and the four-domain mean weights tasks equally. Reported 95% contrast intervals estimate question-sampling uncertainty through a Gaussian parametric bootstrap calibrated on historical paired-question variances. This assumes transfer of the historical calibration, including from a smaller code evaluation; intervals are conditional on the trained models. Training-seed variability remains unquantified. Table 1 collects the endpoint scores and the separate-test means evaluated in Section 6.

3.1 RESPONSE LENGTH CHANGES THE CONTRIBUTION OF EACH DOMAIN AND RESPONSE

Longer responses receive more weight when the loss is averaged over all tokens. In our training batches, mathematics accounts for about three times as many response tokens as instruction following, even though the two domains receive equal numbers of prompts (Figure 1a). Global token averaging therefore gives mathematics a larger share of the loss. Averaging each domain separately fixes its total contribution; averaging response losses equally also removes the extra weight assigned to long responses.




Figure 1: **Loss averaging changes gradient directions; Adam produces more aligned updates and similar final mean scores.** (a) Domain shares under global token averaging (GT). (b) Gradient and update cosines for global token, within-domain token (DT), and within-domain response (DR) averaging; points and bars show batch means and standard deviations. (c–f) Task learning curves under top-64 intersection supervision (I64); S/T mark the initial student/teacher; error bars show model-based question-sampling standard deviations.

The remaining effect of length within a domain has a simple form. For response gradients $g_i = \nabla_{\theta} \ell_i$ and domain means $\bar{g}_d, \bar{T}_d$, the token-weighted gradient is

$$\frac{\sum_{i \in \mathcal{B}_d} T_i g_i}{\sum_{i \in \mathcal{B}_d} T_i} = \bar{g}_d + \frac{\text{Cov}_d(T, g)}{\bar{T}_d}. \quad (5)$$

Token averaging retains the length–gradient covariance term; response averaging uses $\bar{g}_d$. Correcting each domain’s global-token contribution by its target weight divided by its observed token share gives domain-token averaging exactly, as in Open-MOPD (Gao et al., 2026). This correction preserves within-domain length weighting.

3.2 TRAINED ADAM MOMENTS ALIGN UPDATES FROM DIFFERENT AVERAGING RULES

Different gradients produce closely aligned Adam steps. Across six fixed batches at updates 100 and 250 of the sampled-token student, top-64 gradients from equal response weighting and within-domain token weighting have mean cosine 0.68. Each batch contains eight responses per domain and shares student parameters, responses, and teacher scores across averaging rules. Applying those gradients to copies of the same saved Adam state raises their update cosine to 0.96 (Figure 1b). Comparisons with global token averaging show the same pattern.

Adam transforms the current gradient through its accumulated first and second moments (Kingma & Ba, 2015).

In three SmoLLM batches at the exported trained parameters, zero-initialized Adam moments lower the response-versus-token cosine from 0.59–0.90 for raw gradients to 0.41–0.70 for steps. Together, these observations make the optimizer state central to interpreting update alignment.



Figure 2: **Response averaging triggers more clipping; Adam history shapes step norms and one-step loss changes.** (a) Retained gradient fraction and clipping rates for response (DR), within-domain token (DT), and global token (GT) averaging. (b) Step norms after resetting Adam’s first moment or both moments, retaining the step count. (c) Held-out reverse-KL changes (lower is better). Open points are the three batches per checkpoint; filled points and bars show their means and standard deviations. SGD steps match the corresponding Adam step norm.

Clipping preserves gradient direction; optimizer history shapes the step. Equal response weighting triggers clipping throughout training, and token weighting triggers it less often (Figure 2a). Global norm clipping preserves pairwise gradient cosines. Resetting Adam’s moments changes step norms and held-out loss changes, while norm-matched SGD steps separate the averaging rules more strongly (panels b,c). These interventions show how accumulated moments shape the effect of the current gradient.

One-step improvements also vary by domain. We score one held-out response per domain in each batch after BF16 writeback. At update 250, global token averaging changes mean full-vocabulary KL by -0.027×10^{-3} and instruction-following KL by $+0.272 \times 10^{-3}$. The aggregate improvement coexists with a domain-level regression.

3.3 FINAL SCORES ARE CLOSE UNDER ADAM, WITH DIFFERENT TASK TRAJECTORIES

Top-64 supervision yields similar final means across averaging rules. The three rules finish within 0.17 percentage points in mean score. Their learning curves and task rankings vary (Figure 1c–f): all three approach teacher performance on mathematics and instruction following, retain a code gap, and start near the science teacher’s score. Domain token averaging exceeds global token averaging by 0.17 mean points, with an estimated 95% interval of $[-1.27, 1.61]$.

The observed score spread is larger with SGD. With sampled-token supervision, final mean scores span about 1.4 points under SGD and are nearly tied under Adam. Equal response weighting has the highest observed SGD mean. Its advantage over global token weighting is 1.39 points, with an estimated 95% interval of $[-0.04, 2.82]$. Across the evaluated choices, the initial four-domain mean rises from about 34% to 39–40% (Table 1).

4 PARAMETER-UPDATE SPARSITY AND TEACHER ALIGNMENT

We examine numerical precision and shared optimizer state when interpreting concentrated checkpoint changes (Yu et al., 2026) and teacher-update alignment.

4.1 BF16 ROUNDING HIDES WIDESPREAD FP32 CHANGES

Joint training changes about 97% of FP32 master weights. After BF16 rounding, only 7–9% of weights differ visibly from initialization under sampled-token or top-64 supervision with equal response weighting (Figure 3a). Small updates are widespread in FP32 and disappear when the weights are stored in BF16.




Figure 3: **BF16 rounding hides widespread parameter changes, and shared Adam moments align teacher-specific steps.** (a) Fractions of changed coordinates (NZ) and coordinates carrying 90% of squared change (E90), measured in FP32 and BF16. (b,c) Update cosines on domain-specific/shared inputs; $-U_0$ subtracts the zero-gradient step and $m = 0$ resets the first moment. Small/large points show batch/overall means. S/M denote mathematics-only/joint training and PG/I64 denote sampled-token/top-64 supervision. M-DR/DT/GT use top-64 with response/within-domain token/global token averaging; other rows use response averaging.

The precision effect also extends to the concentration of update magnitude. About one third of FP32 coordinates carry 90% of squared cumulative change, compared with about 4% of BF16 coordinates. We count unique model coordinates whose displacement from initialization is nonzero, and rank squared displacements to find the smallest coordinate fraction carrying 90% of the total. Both statistics use the update-500 FP32 master weights and their BF16 stored counterparts.

Loss averaging provides a further comparison between parameter changes and performance. Under top-64 supervision, the averaging rules produce different fractions of changed BF16 weights while reaching similar final mean scores. The changed fraction describes the stored displacement, and its relationship to task performance depends on the training setting.

4.2 SHARED ADAM MOMENTS ALIGN TEACHER-SPECIFIC UPDATES

Teacher-specific steps are strongly aligned at the update-500 sampled-token student with saved Adam state. We compare all six teacher pairs within each of four response batches, averaging pairs within a batch before averaging batches. Shared inputs hold prefixes fixed across teachers; domain-specific inputs use each teacher’s own domain. On responses from each teacher’s own domain, mean update cosines exceed 0.83 under both sampled-token and top-64 supervision. Resetting the first moment while retaining the second, or replacing Adam with norm-matched SGD, reduces the mean cosine to near zero (Figure 3b). The accumulated first moment accounts for much of the alignment.

Adam’s first moment can move parameters even with zero current gradient. To measure the change in a step induced by the current teacher gradient, we subtract this baseline. For update function $U(g; H)$ and saved state H ,

$$\delta U_d = U(g_d; H) - U(0; H). \quad (6)$$

Both steps start from identical parameters and state, so δU_d measures the current gradient’s contribution conditional on the saved moments.

After this subtraction, the mean cosine between teacher-specific step differences falls below 0.01 on domain-specific inputs. When all teachers score the same prefixes, the residual cosine is larger, about 0.24–0.35 (Figure 3c). Teacher-specific alignment therefore depends on both shared optimizer history and shared inputs: removing the zero-gradient component reveals much greater directional separation across domains.




Figure 4: **Top-64 supervision closely approximates the full-vocabulary gradient in Qwen and produces task-dependent learning gains.** (a) Fraction of changed BF16 weights for each local loss at fixed checkpoints. (b) Gradient and update cosines to full-vocabulary supervision. (c) Score differences from sampled-token supervision over training. PG: sampled token; I16/I64: student-teacher top-16/top-64 intersection; T64: teacher top-64. S/M denote mathematics-only/joint training; checkpoint numbers are training updates.

5 VOCABULARY SUPERVISION: PARAMETER UPDATES AND TASK PERFORMANCE

Expanding supervision from one sampled token to a larger vocabulary set changes gradient approximation and task-level learning. We compare sampled-token supervision with the student-teacher top-64 intersection during mathematics-only and joint training. Within each setting, initialization, teachers, response averaging, optimizer settings, and prompt schedule are shared. Mathematics-only training uses four times as many mathematics prompts per update as joint training. Full-vocabulary supervision supplies a reference for the gradient and single-step comparisons.

5.1 BROADER SUPERVISION PRODUCES TASK-DEPENDENT GAINS

The mathematics advantage develops during training. In mathematics-only training, sampled-token supervision leads top-64 by 2.8 accuracy points at update 100; top-64 leads by 2.2 points at update 500 (Figure 4c). Joint training also ends with a mathematics gain from top-64 supervision, accompanied by a science loss: final differences are +2.6 points on mathematics and -1.3 on GPQA. Code and instruction-following differences are smaller.

Top-64’s joint-training mean advantage narrows from 1.69 points at update 250 to 0.43 at update 500, with an estimated endpoint 95% interval of $[-1.01, 1.86]$. These trajectories show that the observed benefit depends on the task and the stage of training.

Larger candidate sets redistribute gains across tasks. Top-16 and top-64 supervision reach development means of 38.90% and 39.04%, and separate-test means of 44.56% and 44.46% (Table 1). On the separate tests, top-64 gains 1.96 mathematics points and loses 2.35 code points relative to top-16. Increasing the candidate set changes both vocabulary support and token weights in Equation 3; the resulting gains vary across domains.

5.2 LARGER CANDIDATE SETS IMPROVE GRADIENT APPROXIMATION

Truncating vocabulary support omits probability mass and changes token weights (Anshumann et al., 2025; Huang & Wei, 2026). We measure coverage as the student probability retained by the student-teacher candidate intersection, $c_k(h) = \sum_{v \in S \cap T} p_\theta(v | h)$, and compare gradients on identical inputs.

Top-64 closely matches the full-vocabulary direction in the Qwen measurements. The intersection retains over 99.9% of student probability on average at the four Qwen checkpoints used for local-step comparisons, with a minimum of 96.55% across their 384 prefixes. At these checkpoints,




Figure 5: **Larger candidate sets improve coverage and gradient approximation in SmoLLM.**

(a,b) Mean retained student probability and fraction of inputs below 90% coverage; bands show response-cluster 95% intervals. Dashed lines mark top-64. (c,d) Gradient cosine and relative error against full-vocabulary supervision at the trained student. Groups comprise typical inputs, uncertain next-token predictions, and inputs with lower top-64 coverage; thin lines show batches and bold lines their means. MixSFT/DR500 denote the initialization/trained student. Coverage and teacher forwards use BF16; student gradient computations use FP32 at exported BF16 parameter values.

its gradient direction is close to that of full-vocabulary supervision (Figure 4b). In an additional fixed-prefix sampling control, averaging 1, 16, and 64 token samples raises the mean gradient cosine from 0.54 to 0.87 to 0.97; top-64 reaches 0.99999 (Figure 6a). This sampling control identifies finite-sample variation as one source of the directional difference.

Lower-coverage inputs retain appreciable gradient error in SmoLLM. SmoLLM3-3B measurements use the MixSFT initialization and its student after 500 updates of top-64 training with response averaging and three teachers for mathematics, code, and instruction following (Gao et al., 2026). Each model generates 384 responses, equally divided among domains, with up to eight audited positions per response. Coverage weights domains, responses, and positions equally at each level. Mean top-64 coverage is 97.42% initially and 98.56% after training; 8.42% and 3.65% of weighted prefixes retain less than 90% mass (Figure 5a,b). Larger candidate sets reduce this low-coverage tail.

Before comparing gradients, we select typical prefixes, high-entropy prefixes, and prefixes with top-64 coverage below 99%, then divide each group into three batches. Across these nine batches at the trained student, top-64 has 2.16–24.54% relative gradient error against full-vocabulary supervision. Top-4 and top-8 produce larger errors in the evaluated batches (Figure 5d). Here relative error is the norm of the gradient difference divided by the full-vocabulary gradient norm, capturing both direction and magnitude. Larger candidate sets improve the approximation, with residual error at top-64.

SmoLLM capability scores also reveal changes below the aggregate. Under shared evaluation prompts and decoding, sampled-token training with SGD raises MATH-500 accuracy from 50.60% to 53.80% (paired question-bootstrap 95% interval for the gain: $[-0.20, 6.60]$). On AIME 2025, the initialization, the top-64/domain-token Adam student, and the sampled-token/response-averaged SGD student each solve four of 30 questions with eight samples per question; only one solved question is common to all three. This comparison describes exported students whose training choices vary together, with incomplete original training records. The fixed-input comparisons above isolate the local effect of vocabulary support within each evaluated student.

5.3 VOCABULARY SUPERVISION LEAVES SINGLE-STEP CONCENTRATION NEARLY UNCHANGED

We compare local steps at updates 100 and 500 of the sampled-token and top-64 students, using 96 prefixes in four batches per checkpoint. At fixed parameters, prefixes, and Adam state, sampled-token, top-64, teacher-top-64, and full-vocabulary losses produce nearly identical single-step con-




Figure 6: **Averaging more token samples improves gradient alignment; Adam history contributes strongly to BF16 changes.** (a) Gradient cosine to full-vocabulary supervision, using 12 fixed prefix subsets with two independent sampling repetitions per budget; I64 uses the top-64 intersection. (b,c) Fractions of changed BF16 weights with zero current gradient or reset moments. PG denotes sampled-token supervision; reset- m keeps the second moment and step count, and fresh Adam resets both moments and the count. Points show batches and bars show means.

centration (Figure 4a). At update 100 of the sampled-token student, all four change about 88% of FP32 coordinates and about 0.15% of BF16 coordinates. This similarity across vocabulary supports is consistent with concentrated stored-weight changes under dense distillation (Yu et al., 2026).

The teacher-top-64 reference restricts the full-KL gradient sum to the teacher’s top-64 tokens using full-vocabulary probabilities. Shared optimizer history contributes strongly to the similar concentration. Even a zero current gradient changes about 0.14% of BF16 parameters in the same comparison, close to the fraction changed by the supervised losses. Resetting Adam’s moments changes the amount of BF16 writeback (Figure 6b,c).

Over training, student parameters, responses, and optimizer moments evolve together. Top-64 and sampled-token supervision consequently accumulate different BF16 displacements (Section 4), alongside the task-dependent learning trajectories above.

6 TASK-LEVEL MONITORING AND CHECKPOINT SELECTION

Separate test questions assess transfer beyond development rankings. We fix 256 additional questions per domain before inference, using Skywork mathematics and TACO code pools, Nemotron instruction-following data, and the natural-science portion of Nemotron knowledge validation. Exact and lexical overlap with the training and known evaluation pools is removed. Scoring uses answer verification, unit tests, instruction constraints, and multiple-choice accuracy, with the same decoding scheme as development evaluation. The changed question distributions and code/instruction scorers define a separate assessment. Table 1 reports its means for all ten final students.

Intermediate checkpoints reveal temporary domain regressions. With sampled-token supervision and response averaging under Adam, science accuracy changes from 32.20% initially to 28.03% at update 50, 30.18% at update 250, and 35.97% at update 500. Tracking individual domains captures this temporary regression and its recovery.

A domain-retention constraint selects the same checkpoints as the mean score. We evaluate checkpoints at updates 50/100/250/500 for Adam and 100/250/500 for SGD. We compare the highest development mean with the highest mean subject to

$$s_d(\theta) \geq s_d(\theta_0) - \epsilon_d \quad \text{for all } d, \quad \epsilon_d = 2 \text{ percentage points}, \quad (7)$$

where s_d is the domain score and θ_0 the initial student. Both rules retain the initialization unless the mean improves by at least one point; choices are fixed before separate-test inference.

Both rules select update 500 in nine trajectories and update 250 for sampled-token training with domain-token averaging under Adam. The earlier checkpoint scores 43.14% on the separate test,

compared with 43.51% at update 500. The retention constraint is inactive at every selected checkpoint, while the learning curves contain temporary domain regressions.

Method rankings vary with evaluation questions. For sampled-token supervision with response averaging, SGD leads Adam by 1.35 mean points on development benchmarks; separate-test means are 43.48% and 43.55%, respectively. A further sensitivity check removes four MATH-500 questions with known training-pool matches: four-domain means change by 0.09–0.16 points, reordering some near-tied methods while preserving the highest observed mean for response-averaged SGD. Together, these results support reading aggregate rankings alongside per-task trajectories and the evaluation questions that produced them.

7 RELATED WORK

Distillation objectives and vocabulary truncation. MiniLLM uses reverse KL (Gu et al., 2024); Generalized Knowledge Distillation (GKD) combines student-generated sequences with alternative divergences (Agarwal et al., 2024); f -DISTILL derives word-level losses from sequence-level f -divergences (Wen et al., 2023). Sparse Logit Sampling analyzes top- k caching bias (Anshumann et al., 2025), while Tail-Aware OPD restores tail-mass information lost by top- k renormalization (Huang & Wei, 2026). Our intersection loss retains full-softmax log ratios with detached student-top- k weights. Section 5 compares its coverage, gradient error, updates, and task scores with sampled-token and full-vocabulary supervision.

Multi-teacher distillation and loss normalization. MOPD integrates specialized teachers on student-generated responses (Ma et al., 2026); Open-MOPD adds domain token balancing and adaptive weighting (Gao et al., 2026). DAPO uses token-level averaging (Yu et al., 2025), and Dr. GRPO analyzes length and reward-normalization bias (Liu et al., 2025). We separate domain weighting from within-domain response weighting: Open-MOPD’s token-share correction gives DT for fixed batches and domain weights, while DR also removes response-length weighting.

Post-training parameter geometry. Mukherjee et al. (2025) study small parameter subsets supporting RL fine-tuning; Yu et al. (2026) report concentrated distillation updates and overlap between separately trained students. Shen et al. (2026) examine OPD geometry and response-position selection. We compare teachers at the same student and Adam state, distinguish FP32 from BF16 changes, and vary vocabulary support at fixed prefixes.

8 DISCUSSION

The path from a distillation loss to task performance passes through optimizer history, numerical precision, and an evolving training trajectory. In our comparisons, Adam aligns directionally different gradients, and BF16 rounding hides widespread FP32 changes. Broader vocabulary supervision improves gradient approximation and produces task-dependent learning gains, including a mathematics–science trade-off and changing advantages over time. These findings give complementary roles to the measurements: gradients describe the immediate training signal, parameter changes show its interaction with the optimizer and precision, and per-task learning curves reveal the resulting capability gains and regressions.

REPRODUCIBILITY STATEMENT

The evidence bundle includes figure and table values, configurations, source hashes, scoring audits, and plotting scripts. The appendix provides full protocols, supporting analyses, and overlap checks.

ETHICS STATEMENT

We use public models and benchmarks. Generative AI assisted manuscript editing and analysis code; the authors are responsible for the measurements and their interpretation.

REFERENCES

- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2306.13649>.
- Anshumann, Mohd Abbas Zaidi, Akhil Kedia, Jinwoo Ahn, Taehwak Kwon, Kangwook Lee, Haejun Lee, and Joohyung Lee. Sparse logit sampling: Accelerating knowledge distillation in LLMs. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 18085–18108. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.885. URL <https://aclanthology.org/2025.acl-long.885/>.
- DeepSeek-AI, Anyi Xu, Bangcai Lin, Bing Xue, Bingxuan Wang, Bingzheng Xu, Bochao Wu, Bowei Zhang, Chaofan Lin, Chen Dong, Chenchen Ling, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chengyu Hou, Chenhao Xu, Chenze Shao, Chong Ruan, Conner Sun, Damai Dai, Daya Guo, Dejian Yang, Deli Chen, Donghao Li, Dongjie Ji, Erhang Li, Fang Wei, Fangyun Lin, Fangzhou Yuan, Feiyu Xia, Fucong Dai, Guangbo Hao, Guanting Chen, Guoai Cao, Guolai Meng, Guowei Li, Han Yu, Han Zhang, Hanwei Xu, Hao Li, Haofen Liang, Haoling Zhang, Haoming Luo, Haoran Wei, Haotian Yuan, Haowei Zhang, Haowen Luo, Haoyu Chen, Haozhe Ji, Hengqing Zhang, Honghui Ding, Hongxuan Tang, Huanqi Cao, Huazuo Gao, Hui Qu, Hui Zeng, J Yang, JQ Zhu, Jia Luo, Jia Song, Jia Yu, Jialiang Huang, Jialu Cai, Jian Liang, Jiangting Zhou, Jiasheng Ye, Jiashi Li, Jiaxin Xu, Jiewen Hu, Jieyu Yang, Jin Chen, Jin Yan, Jingchang Chen, Jingli Zhou, Jingting Xiang, Jingyang Yuan, Jingyuan Cheng, Jingzi Zhou, Jinhua Zhu, Jiping Yu, Joseph Sun, Jun Ran, Junguang Jiang, Junjie Qiu, Junlong Li, Junmin Zheng, Junxiao Song, Kai Dong, Kaige Gao, Kang Guan, Kexing Zhou, Kezhao Huang, Kuai Yu, Lean Wang, Lecong Zhang, Lei Wang, Leyi Xia, Li Zhang, Liang Zhao, Lihua Guo, Lingxiao Luo, Linwang Ma, Linyan Zhu, Litong Wang, Liyu Cai, Liyue Zhang, Longhao Chen, MS Di, MY Xu, Max Mei, Miaojuan Wang, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Mingming Li, Mingxu Zhou, Minmin Han, Ning Wang, Panpan Huang, Panpan Wang, Peixin Cong, Peiyi Wang, Peng Zhang, Qiancheng Wang, Qihao Zhu, Qingyang Li, Qinyu Chen, Qiushi Du, Qiwei Jiang, Rui Tian, Ruifan Xu, Ruijie Lu, Ruiling Xu, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, Runqian Chen, Runqiu Yin, Runxin Xu, Ruomeng Shen, Ruoyu Zhang, Ruyi Chen, SH Liu, Shanghao Lu, Shangmian Sun, Shangyan Zhou, Shanhuang Chen, Shaofei Cai, Shaocheng Nie, Shaoqing Wu, Shaoyuan Chen, Shengding Hu, Shengyu Liu, Shiqiang Hu, Shirong Ma, Shiyu Wang, Shuiping Yu, Shunfeng Zhou, Shuting Pan, Shuying Yu, Songyang Zhou, Tao Ni, Tao Yun, Tian Jin, Tian Pei, Tian Ye, Tianle Lin, Tianran Ji, Tianyi Cui, Tianyuan Yue, Tingting Yu, Tun Wang, W Zhang, WL Xiao, Wangding Zeng, Wei An, Weilin Zhao, Wen Liu, Wenfeng Liang, Wenjie Pang, Wenjing Luo, Wenjing Yao, Wenjun Gao, Wenkai Yang, Wenlve Huang, Wenqing Hou, Wentao Zhang, Wenting Ma, Xi Gao, Xiang He, Xiangwen Wang, Xianzu Wang, Xiao Bi, Xiaodong Liu, Xiaohan Wang, Xiaokang Chen, Xiaokang Zhang, Xiaotao Nie, Xiaowen Sun, Xiaoxiang Wang, Xin Cheng, Xin Liu, Xin Xie, Xingchao Liu, Xingchen Liu, Xingkai Yu, Xingyou Li, Xinyu Yang, Xinyu Zhang, Xu Chen, Xuanyu Wang, Xuecheng Su, Xueyin Chen, Xuheng Lin, Xuwei Fu, YC Yan, YQ Wang, YW Ma, Yanfeng Luo, Yang Zhang, Yanhong Xu, Yanru Ma, Yanwen Huang, Yao Li, Yao Li, Yao Xu, Yao Zhao, Yaofeng Sun, Yaohui Wang, Yi Qian, Yi Shao, Yi Yu, Yichao Zhang, Yifan Ding, Yifan Shi, Yijia Wu, Yiliang Xiong, Yiling Ma, Ying He, Ying Tang, Ying Zhou, Yingjia Luo, Yinmin Zhong, Yishi Piao, Yisong Wang, Yixiang Zhang, Yixiao Chen, Yixuan Tan, Yixuan Wei, Yiyang Ma, Yiyuan Liu, Yonglun Yang, Yongqiang Guo, Yongtong Wu, Yu Wu, YuKun Li, Yuan Cheng, Yuan Ou, Yuanfan Xu, Yuanhao Li, Yudian Wang, Yuehan Yang, Yuer Xu, Yuhan Wu, Yuhao Meng, Yuheng Zou, Yukun Zha, Yunfan Xiong, Yupeng Chen, Yuping Lin, Yuqian Cao, Yuqian Wang, Yushun Zhang, Yuting Yan, Yutong Lin, Yuxian Gu, Yuxiang Luo, Yuxiang You, Yuxuan Liu, Yuxuan Zhou, Yuyang Zhou, Yuzhen Huang, ZF Wu, Zehao Wang, Zehua Zhao, Zehui Ren, Zekai Zhang, Zhangli Sha, Zhe Fu, Zhe Ju, Zhean Xu, Zhenda Xie, Zhengyan Zhang, Zheren Gao, Zhewen Hao, Zhibin Gou, Zhicheng Ma, Zhigang Yan, Zhihong Shao, Zhixian Huang, Zhixuan Chen, Zhiyu Wu, Zhizhou Ren, Zhongyu Wu, Zhuoshu Li, Zhuping Zhang, Zian Xu, Zihao Wang, Zihua Qu, Zihui Gu, Zijia Zhu, Zilin Li, Zipeng Zhang, Ziwei Xie, Ziyi Gao, Ziyi Wan, Zizheng Pan, and Zongqing Yao. Deepseek-v4: Towards highly efficient million-token context intelligence, 2026. URL <https://arxiv.org/abs/2606.19348>.

- Huan-ang Gao, Haohan Chi, Yong Yan, Shiyuan Feng, Hanlin Wu, Zheng Jiang, Bingxiang He, Wei-Ying Ma, Ya-Qin Zhang, and Hao Zhou. Open-MOPD: Diagnosing and fixing capability imbalance in multi-teacher on-policy distillation, 2026. URL <https://arxiv.org/abs/2608.19098v1>.
- Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large language models. In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2306.08543v4>.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems*, 2021. doi: 10.48550/arXiv.2103.03874. URL <https://arxiv.org/abs/2103.03874v2>.
- Huipeng Huang and Hongxin Wei. Tail-aware top- k on-policy distillation, 2026. URL <https://arxiv.org/abs/2608.14728>.
- Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. Livecodebench: Holistic and contamination free evaluation of large language models for code, 2024. URL <https://arxiv.org/abs/2403.07974v2>.
- Kimi Team, Tongtong Bai, Yifan Bai, Yiping Bao, M. C., Jianfeng Cai, Xinyuan Cai, Peizhou Cao, Yuxuan Cao, Ziwei Chai, Y. Charles, H. S. Che, Guanduo Chen, Guangyu Chen, Guanzheng Chen, Huarong Chen, Jia Chen, Jianlong Chen, Jun Chen, Kexin Chen, Peng Chen, Ruijue Chen, Wentao Chen, Xin Chen, Yang Chen, Yanru Chen, Yifei Chen, Yingjiang Chen, Yuankun Chen, Yujie Chen, Yutian Chen, Zhirong Chen, Dazhi Cheng, Yean Cheng, Jialei Cui, Jingbing Cui, Anqi Dai, Jiaqi Deng, Hao Ding, Rui Ding, Shaofeng Ding, Mengfan Dong, Mengnan Dong, Yuhao Dong, Yuxin Dong, Angang Du, Chenzhuang Du, Dikang Du, Jusen Du, Yulun Du, Yu Fan, Jing Feng, Qiulin Feng, Yichen Feng, Kelin Fu, Qiang Fu, Fuxuan Gao, Hongcheng Gao, Jingyue Gao, Tong Gao, Weijia Gao, Shangyi Geng, Jie Gong, Linhu Gong, Shengao Gong, Xiaochen Gong, Qizheng Gu, Yicheng Gu, Shuhao Guan, Haiqing Guo, Shiqi Guo, Xiang Guo, Zhengyan Guo, Beixi Hao, Wenxin Hao, Xiaoru Hao, Dailan He, Haotian He, Lehan He, Qi He, Weiran He, Xinran He, Xinyi He, Yibo He, Yunjia He, Chao Hong, Tiange Hong, Hao Hu, Jiaxi Hu, Ruikun Hu, Weiming Hu, Yangyang Hu, Zhenxing Hu, Liang Hua, Jinbin Huang, Ke Huang, Ruiyuan Huang, Siying Huang, Weixiao Huang, Yan Huang, Zhengjie Huang, Zhiqi Huang, Yulong Hui, Chaobo Jia, Yutong Jiang, Zhejun Jiang, Zuoyou Jiang, Wenyi Jin, Xinyi Jin, Yu Jing, Huanjun Kong, Guokun Lai, Aidi Li, Cheng Li, Chengyuan Li, Cong Li, Fang Li, Guanyu Li, Haoyang Li, Jia Li, Junxiong Li, Lei Li, Letian Li, Lincan Li, Weihong Li, Wentao Li, Xintong Li, Yang Li, Yishen Li, Yiwei Li, Yuxiao Li, Zhaowei Li, Zhaoxi Li, Zheming Li, Zhengxiao Li, Zhiyuan Li, Jiawei Lin, Xiaohan Lin, Yibo Lin, Zichao Lin, Ziyang Lin, Bill Liu, Boxiao Liu, Chuan Liu, Liang Liu, Shaowei Liu, Shudong Liu, Shuran Liu, Tianwei Liu, Weizhou Liu, Yangyang Liu, Yanming Liu, Yibo Liu, Yipeng Liu, Zhengying Liu, Zhiheng Liu, Enzhe Lu, Haoyu Lu, Lin-qiang Lu, Tingzhan Lu, Zhiyuan Lu, Aotian Luo, G. Luo, Junyu Luo, Yifan Luo, B. Lyu, Wenzhou Lyu, Shaoguang Mao, Yuan Mei, Xin Men, Mingqing Ni, Yixuan Niu, Siyuan Pan, Shujun Peng, Zhangyang Qi, Ruoyu Qin, ZeChao Qin, Zeyu Qin, Haiquan Qiu, Jianxin Qiu, Jiezhong Qiu, Bowen Qu, Yuhao Qu, Zeyu Shang, Youbo Shao, Han Shen, Jincheng Shi, Juanfeng Shi, Lidong Shi, Shengyuan Shi, Wingchun Siu, Pengwei Song, Xiaoxi Song, Jianlin Su, Yunfeng Su, Zhaochen Su, Lin Sui, Jingsong Sun, Junyao Sun, Shaoning Sun, Shuzhe Sun, Tongyu Sun, Yujun Sun, Yunpeng Tai, Chuning Tang, Heyi Tang, Sirui Tang, Zecheng Tang, Chaoran Tian, Rongpeng Tian, Yu Tian, Wei Tu, Chensi Wang, Chuang Wang, Chunjie Wang, Dinglu Wang, Feng Wang, Hailong Wang, Haiming Wang, Hao Wang, Hao Wang, Huaqing Wang, Hui Wang, Jiayi Wang, Jinglong Wang, Jinhong Wang, Jiuzheng Wang, Linian Wang, Shaobo Wang, Shen-zhi Wang, Shuyi Wang, Si Wang, Siyuan Wang, Tianfu Wang, Wenjue Wang, Xingran Wang, Xinmei Wang, Xinyuan Wang, Xusheng Wang, Yalin Wang, Yangkun Wang, Yao Wang, Yaoyu Wang, Yejie Wang, Yiqin Wang, Yucheng Wang, Yuzhi Wang, Zhaoji Wang, Zhaowei Wang, Zhengtao Wang, Zhenhao Wang, Zhongsheng Wang, Zifan Wang, Chu Wei, Ming Wei, Shouxin Wei, Zichen Wen, Fan Wu, Haoning Wu, Rucong Wu, Wenhao Wu, Xiaoxue Wu, Yingcong Wu, Yongqi Wu, Yuxin Wu, Zijian Wu, Xinglang Xian, Chenxuan Xiang, Yuye Xiang, Bocheng Xiao, Chenjun Xiao, Xin Xiao, Jin Xie, Xiaotong Xie, Yifeng Xie, Zhe Xie, Bowei Xing, Yiming

- Xiong, Baosheng Xu, Boyu Xu, Jiale Xu, Jianfan Xu, Jing Xu, Jinjing Xu, L. H. Xu, Qingtao Xu, Shuyao Xu, Suting Xu, Tiantian Xu, Tianxiang Xu, Weixin Xu, Xinran Xu, Yangchuan Xu, Ye Xu, Yueni Xu, Ziyao Xu, Haonan Xue, Junjie Yan, Yaoyao Yan, Fan Yang, Guangyao Yang, Hao Yang, Junwei Yang, Ruoyu Yang, Wenjie Yang, Xiaofei Yang, Xinyu Yang, Yi Yang, Yiling Yang, Ying Yang, Yuchen Yang, Zhen Yang, Zhilin Yang, Zian Yang, Zuhao Yang, Haotian Yao, Dan Ye, Haoran Ye, Wenjie Ye, Zhanbo Ye, Bohong Yin, Haoxiang Yin, Xietong Yin, Chengzhen Yu, Haozhen Yu, Longhui Yu, Shengnan Yu, Shuying Yu, Tianxiang Yu, Enming Yuan, Mengjie Yuan, Tongtian Yue, Wei Yue, Yang Yue, Dunyuan Zha, Haobing Zhan, B. H. Zhang, Dehao Zhang, Fei Zhang, Hao Zhang, Haoyuan Zhang, Huanyu Zhang, Jiawei Zhang, Jiaxuan Zhang, Jin Zhang, Kaiyi Zhang, Miaozhen Zhang, Puqi Zhang, Qinglei Zhang, Rong Zhang, Rui Zhang, Shaoshuai Zhang, Shiyi Zhang, Xiaobin Zhang, Xiaoyun Zhang, Y. Zhang, Yangkun Zhang, Ye Zhang, Yichi Zhang, Yikun Zhang, Yizhi Zhang, Yongting Zhang, Yu Zhang, Yutao Zhang, Yutong Zhang, Zheng Zhang, Zijing Zhang, Bin Zhao, Chenguang Zhao, Feifan Zhao, Jinglun Zhao, Jinxiang Zhao, Shuai Zhao, Wenshuo Zhao, Xiangyu Zhao, Xuanle Zhao, Yikai Zhao, Zijia Zhao, Haozhi Zheng, Huabin Zheng, Ruihan Zheng, Shaojie Zheng, Tengyang Zheng, Haofeng Zhong, Lei Zhong, Longguang Zhong, M. Zhou, Qianqiang Zhou, Runjie Zhou, Ruozhang Zhou, Xinyu Zhou, Yiqiao Zhou, Zaida Zhou, Jinguo Zhu, Liya Zhu, Xinhao Zhu, Yangjunfeng Zhu, Yuxuan Zhu, Zhen Zhu, Chen Zhuang, Weiyu Zhuang, and Xinxing Zu. *Kimi k3: Open frontier intelligence*, 2026. URL <https://arxiv.org/abs/2607.24653>.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations*, 2015. URL <https://arxiv.org/abs/1412.6980>.
- Zichen Liu, Changyu Chen, Wenjun Li, Penghui Qi, Tianyu Pang, Chao Du, Wee Sun Lee, and Min Lin. Understanding R1-Zero-like training: A critical perspective, 2025. URL <https://arxiv.org/abs/2503.20783>.
- Wenhan Ma, Jianyu Wei, Liang Zhao, Hailin Zhang, Bangjun Xiao, Lei Li, Qibin Yang, Bofei Gao, Yudong Wang, Rang Li, Jinhao Dong, Zhifang Sui, and Fuli Luo. MOPD: Multi-teacher on-policy distillation for capability integration in llm post-training, 2026. URL <https://arxiv.org/abs/2606.30406v1>.
- Sagnik Mukherjee, Lifan Yuan, Dilek Hakkani-Tur, and Hao Peng. Reinforcement learning finetunes small subnetworks in large language models, 2025. URL <https://arxiv.org/abs/2505.11711v2>. Version 2; NeurIPS 2025.
- NVIDIA. Nemotron 3 Ultra: Open, efficient mixture-of-experts hybrid Mamba-Transformer model for agentic reasoning, 2026. URL <https://arxiv.org/abs/2606.15007v1>.
- Valentina Pyatkin, Saumya Malik, Victoria Graf, Hamish Ivison, Shengyi Huang, Pradeep Dasigi, Nathan Lambert, and Hannaneh Hajishirzi. Generalizing verifiable instruction following, 2025. URL <https://arxiv.org/abs/2507.02833v3>.
- David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. GPQA: A graduate-level google-proof q&a benchmark, 2023. URL <https://arxiv.org/abs/2311.12022v1>.
- Zhennan Shen, Yanshu Li, Qingyu Yin, Chak Tou Leong, Zhilin Wang, Yanxu Chen, Rongduo Han, Sunbowen Lee, and Yi R. Fung. On the geometry of on-policy distillation, 2026. URL <https://arxiv.org/abs/2606.07082v3>. Version 3, June 14, 2026.
- Yuqiao Wen, Zichao Li, Wenyu Du, and Lili Mou. f -divergence minimization for sequence-level knowledge distillation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 10817–10834. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.acl-long.605. URL <https://aclanthology.org/2023.acl-long.605/>.
- Xiaomi LLM-Core Team. MiMo-V2-Flash technical report, 2026. URL <https://arxiv.org/abs/2601.02780v2>.

- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jing Zhou, Jingren Zhou, Junyang Lin, Kai Dang, Keqin Bao, Kexin Yang, Le Yu, Lianghao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang Ren, Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yinger Zhang, Yu Wan, Yuqiong Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan Qiu. Qwen3 technical report, 2025. URL <https://arxiv.org/abs/2505.09388v1>.
- Guo Yu, Wenlin Liu, Yulan Hu, Hao-Xuan Ma, Jun-Peng Jiang, and Han-Jia Ye. Dense supervision, sparse updates: On the sparsity and geometry of on-policy distillation, 2026. URL <https://arxiv.org/abs/2606.13657v3>. Version 3, July 30, 2026.
- Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou, Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, and Mingxuan Wang. DAPO: An open-source LLM reinforcement learning system at scale, 2025. URL <https://arxiv.org/abs/2503.14476>.

Table 2: **Task gains coexist with domain-level trade-offs.** Joint-training endpoint scores (%); Mean weights domains equally. PG: sampled-token supervision; I16/I64: top-16/top-64 intersection supervision. DR/DT: equal domain weights with response/token averaging; GT: global token averaging. Δ is the mean difference from PG-DR/Adam (percentage points). Brackets give model-based paired 95% intervals (Appendix A.3).

Configuration	MATH-500	Code	IF	GPQA	Mean	Δ [95% interval]
Initial student	68.80	16.41	18.33	32.20	33.94	-4.68 [-6.42, -2.93]
<i>Adam</i>						
PG-DR	73.20	18.27	27.00	35.97	38.61	Reference
PG-DT	73.20	19.24	27.00	34.96	38.60	-0.01 [-1.45, 1.43]
PG-GT	76.20	17.92	27.33	32.94	38.60	-0.01 [-1.44, 1.42]
I16-DR	75.40	17.75	27.00	35.46	38.90	+0.29 [-1.15, 1.73]
I64-DR	75.80	18.30	27.33	34.71	39.04	+0.43 [-1.01, 1.86]
I64-DT	74.00	18.86	28.67	35.22	39.19	+0.58 [-0.87, 2.02]
I64-GT	74.20	19.80	28.67	33.41	39.02	+0.41 [-1.03, 1.85]
<i>SGD</i>						
PG-DR	76.40	19.53	28.33	35.59	39.96	+1.35 [-0.09, 2.79]
PG-DT	75.20	20.31	28.33	33.90	39.44	+0.83 [-0.61, 2.26]
PG-GT	75.00	19.26	25.33	34.71	38.58	-0.04 [-1.47, 1.40]

A EXPERIMENTAL PROTOCOL AND RESULTS

A.1 MODEL, DATA, AND TRAINING PROTOCOL

The main-study student is Qwen3-1.7B (Yang et al., 2025); training uses seeds 42, 43, and 44. The four teachers were trained by reinforcement learning from the same initial model, specializing in mathematics, code, instruction following, and science, respectively. Student and teachers share the tokenizer and native chat template with thinking disabled. The experiment records identify the model revisions and weight files.

AdamW configurations use learning rate 2.5×10^{-7} with $(\beta_1, \beta_2) = (0.9, 0.98)$, $\epsilon = 10^{-8}$, no weight decay or schedule, and global norm-1 clipping. Training sampling uses temperature 1 and top- $p = 1$, with a 2,048-token prompt cap and 4,096-token response cap. One fresh response is sampled per prompt. Each optimizer update aggregates 64 responses; domain allocation is specified below.

The multi-teacher setup allocates 16 prompts to each domain per update and uses target domain weights 1/4. The single-teacher setup allocates all 64 prompts to the same domain. I64 uses the top-64 objective in Equation 3; I16 uses $k = 16$.

A.2 EVALUATION AND SCORING

Each main-study capability suite contains MATH-500 (500 questions) (Hendrycks et al., 2021), LiveCodeBench v6 (1,055 questions) (Jain et al., 2024), IFBench strict (300 questions) (Pyatkin et al., 2025), and GPQA Diamond (198 questions) (Rein et al., 2023). Mathematics, code, and instruction following use one greedy response per question. GPQA uses four responses per question at temperature 1 and top- $p = 1$. Its score is the mean correctness over all four responses, then over questions: average@4 accuracy. Evaluation permits up to 32,768 response tokens, bounded by the remaining configured 32,768-token context. Each complete suite contains 2,053 questions and 2,647 responses per model evaluation.

The expanded development grid comprises updates of 50/100/250/500 for seven Adam trajectories and 100/250/500 for three SGD trajectories, plus the initial student evaluation. Appendix G.1 describes additional test sets for the same four domains. Table 2 gives the complete joint-training endpoint scores. Table 3 adds single-teacher scores and intermediate checkpoints.

Table 3: **PG/I64 performance depends on the task and checkpoint.** Mean scores (%): S is mathematics-only training and M is joint training; both use DR. Joint endpoints appear in Table 2.

Configuration	Updates	Math	Code	IF	GPQA
S-PG	100	77.40	18.75	18.67	31.82
S-PG	250	73.80	20.31	17.00	35.35
S-PG	500	73.80	19.53	18.00	31.82
S-I64	100	74.60	18.75	17.67	33.21
S-I64	250	74.40	21.09	18.33	32.45
S-I64	500	76.00	17.97	18.67	30.05
M-PG	50	69.60	16.41	18.67	28.03
M-PG	100	72.00	17.19	19.33	31.19
M-PG	250	73.40	16.41	26.00	30.18
M-I64-DR	50	69.60	17.19	19.00	31.69
M-I64-DR	100	71.60	17.97	19.00	32.70
M-I64-DR	250	74.80	17.97	26.00	33.96

Table 4: **Reference scores for the four domain-specific teachers (%)**. Science uses GPQA average@4.

Teacher / target domain	Responses	Score (%)	Truncation (%)
Math	500	75.60	6.80
Code	1055	23.44	24.22
If	300	29.00	7.33
Science	792	36.40	0.51

The four RL teachers were evaluated separately on their respective domain benchmarks using the same prompt and decoding protocol (Table 4). These scores quantify each teacher’s performance on its target domain.

A.3 CAPABILITY UNCERTAINTY

Table 5 reports the model-based sampling intervals cited in the main text. They estimate finite-question uncertainty using historical paired-question variance calibration, conditional on the reported training outcomes. Their interpretation depends on that calibration; training-seed variability requires a separate estimate from replicated trajectories. The calculation uses 100,000 Gaussian parametric-bootstrap replicates (seed 1042), independent domains, and question counts of 500/1,055/300/198 for mathematics/code/instruction following/GPQA. Domain-specific variance ratios κ are calibrated from historical paired-question intervals relative to independent Bernoulli variances, with score-lattice variance floors. Transferring these ratios from the historical 128-question code evaluation to the current 1,055-question protocol is an assumption. GPQA retains 198 question clusters. The table also reports sensitivity checks using $\min(2\kappa, 1)$ and the independent reference $\kappa = 1$.

B LENGTH-WEIGHTING DERIVATION

Fix a rollout batch, its valid-token masks, and domain loss weights λ_d . The PG log-ratio coefficients are treated as constants during backpropagation. For $n_d = |\mathcal{B}_d|$, let $\bar{T}_d = n_d^{-1} \sum_{i \in \mathcal{B}_d} T_i$ and $\bar{g}_d = n_d^{-1} \sum_{i \in \mathcal{B}_d} g_i$. Define the empirical covariance using the same $1/n_d$ normalization:

$$\text{Cov}_d(T, g) = \frac{1}{n_d} \sum_{i \in \mathcal{B}_d} (T_i - \bar{T}_d)(g_i - \bar{g}_d). \quad (8)$$

Expanding gives $n_d^{-1} \sum_i T_i g_i - \bar{T}_d \bar{g}_d$. Dividing by $\bar{T}_d$ gives Equation 5. GT combines the resulting domain gradients using observed token shares; DT combines them using λ_d ; DR combines $\bar{g}_d$ using λ_d . Multiplying each domain’s globally averaged terms by its target weight divided by its observed token share therefore gives DT exactly.

Table 5: **Small endpoint contrasts remain uncertain under the calibrated model.** Mean differences (percentage points) and estimated paired 95% intervals; 2κ weakens the assumed variance reduction, and $\kappa = 1$ uses the independent Bernoulli reference. These are pointwise, model-based intervals.

Contrast	Δ	Calibrated	$\min(2\kappa, 1)$	$\kappa = 1$
I64/Adam: DT – DR	+0.15	[-1.28, 1.59]	[-1.88, 2.18]	[-3.20, 3.51]
I64/Adam: GT – DR	-0.02	[-1.44, 1.41]	[-2.04, 2.01]	[-3.36, 3.33]
I64/Adam: DT – GT	+0.17	[-1.27, 1.61]	[-1.87, 2.20]	[-3.19, 3.53]
PG/Adam: DT – DR	-0.01	[-1.45, 1.43]	[-2.05, 2.03]	[-3.37, 3.35]
PG/Adam: GT – DR	-0.01	[-1.44, 1.42]	[-2.04, 2.01]	[-3.35, 3.33]
PG/Adam: DT – GT	+0.00	[-1.43, 1.43]	[-2.02, 2.02]	[-3.33, 3.34]
PG/SGD: DT – DR	-0.53	[-1.96, 0.90]	[-2.55, 1.49]	[-3.88, 2.83]
PG/SGD: GT – DR	-1.39	[-2.82, 0.04]	[-3.41, 0.63]	[-4.73, 1.96]
PG/SGD: DT – GT	+0.86	[-0.57, 2.29]	[-1.16, 2.88]	[-2.48, 4.20]
PG-DR: SGD – Adam	+1.35	[-0.09, 2.79]	[-0.68, 3.39]	[-2.01, 4.71]
PG-DT: SGD – Adam	+0.84	[-0.60, 2.27]	[-1.20, 2.87]	[-2.52, 4.19]
PG-GT: SGD – Adam	-0.02	[-1.44, 1.40]	[-2.03, 1.99]	[-3.34, 3.30]
Adam/DR: I64 – PG	+0.43	[-1.01, 1.86]	[-1.61, 2.46]	[-2.93, 3.78]
Adam/DR: I64 – I16	+0.13	[-1.30, 1.56]	[-1.89, 2.15]	[-3.21, 3.47]

The derivation is a finite-batch weighted-average identity. Response sampling, masks, and lengths are held fixed. Its population analogue involves an expectation of a ratio; replacing that expectation by a ratio of expectations requires additional assumptions or a large-batch limit.

B.1 FIXED-BATCH GRADIENT COMPARISON

At M-PG checkpoints 100 and 250, three cached response batches (seeds 1042, 1043, and 1044) each contain eight responses per domain. Responses use the training prompt format and 4,096-token response cap. For each batch, we compute each response’s mean I64 gradient once at the saved BF16 model values using FP32 computation, then aggregate those gradients under DR, DT, and GT. All three averaging rules use the same responses, token masks, teacher scores, and model coordinates. Each comparison starts from the same student parameters and Adam first and second moments.

Table 6 reports every pairwise raw-gradient cosine. The domain-wise covariance decomposition is checked in FP64 for all 24 batch-domain groups; the largest relative residual is 1.53×10^{-15} . This numerical check verifies the aggregation identity. The measured cosine differences show that response length and gradient direction are correlated in these batches.

The optimizer interventions in Figure 2 reuse these six batches. Each aggregate gradient is globally clipped at norm 1 and applied to the same saved FP32 Adam state for that batch. The saved-state branch retains both moments; the reset- m branch zeros the first moment while retaining the second moment and optimizer step counter; reset- (m, v) zeros both moments while retaining the step counter. The momentum-free SGD control follows the negative gradient and is rescaled separately for DR, DT, and GT to match the FP32 Adam step norm for that averaging rule. After each FP32 update, the model parameters are cast to BF16 and evaluated on one held-out response per domain. Evaluation uses full-vocabulary reverse KL; the experiment archive also retains the I64 surrogate used for training. Each comparison applies a single optimizer step to a checkpoint copy.

C PARAMETER AND OPTIMIZER MEASUREMENTS

C.1 COORDINATES AND RECORDED QUANTITIES

Training measurements separately record gradients before and after clipping, FP32 optimizer steps, and cumulative parameter changes in the FP32 master weights and BF16 model weights. Cumulative changes are measured from the same initialization. All global statistics use unique model coordinates, excluding vocabulary padding and duplicate entries for tied weights. Per-layer statistics and absolute-threshold sensitivities are retained in the experiment archive. For an absolute threshold τ ,

Table 6: **Response weighting changes gradient directions at equal domain weights.** Raw-gradient cosines on identical cached responses. DR and DT use equal domain loss weights; only DT weights responses by length.

Updates	Batch	DR-DT	DR-GT	DT-GT
100	1	0.7838	0.6219	0.9266
100	2	0.9110	0.5180	0.7010
100	3	0.5780	0.6659	0.9745
250	1	0.6240	0.4072	0.8742
250	2	0.7864	0.3302	0.4953
250	3	0.3990	0.3569	0.9612




Figure 7: **Update sparsity depends on the numerical threshold.** FP32 steps and BF16 parameter changes; lines show batch means and shading their range.

sparsity is the fraction of coordinates with magnitude at most τ ; the figures show the fraction exceeding τ . Energy concentration uses squared Euclidean norm and ranks coordinates by absolute change.

For checkpoint-based measurements, the exported checkpoint contains FP32 master values, BF16 model values, first and second moments, optimizer hyperparameters, and step counters. Student and teacher values are verified against their saved models in the same precision. Forward and backward passes evaluate these values in FP32 with TF32 disabled. All proposed steps are computed from the same saved state with a non-fused Adam implementation. The online training optimizer uses its own fused implementation.

We compare three Adam states at the same checkpoint. *Saved* retains both moments and the step counter. *Reset* zeros the first moment while retaining the second moment and step counter. *Fresh* zeros both moments and resets the step counter. Global norm clipping at 1 is applied before the Adam calculation, and weight decay is zero. The plotted controls use two cached response batches; the zero-gradient control is evaluated once. For the zero-gradient branch, the saved first moment continues to decay and move the master weights. Casting the saved FP32 master weights to BF16 without an optimizer step exactly reproduces the saved model. Figure 7 shows how the threshold changes the measured fraction of updated coordinates.

Setting the current gradient to zero while retaining the trained Adam moments changes 0.1403% of BF16 coordinates, compared with 0.1486% for PG. Resetting the first moment gives 0.0504% for PG; resetting both moments and the step counter gives 0.5140%. The parameters changed by Adam depend jointly on the current gradient, accumulated moments, and BF16 rounding. These controls supplement the update-sparsity measurements in Sections 4 and 5.

Figure 6(b,c) displays these BF16 controls. The full FP32 and BF16 measurements are retained in the experiment archive.




Figure 8: **Gradient overlap depends strongly on the input domain.** Left/center: teacher-by-input-domain controls. Right: top-1% gradient Jaccard versus cosine on domain-specific or shared inputs.

C.2 CACHED RESPONSES, GRADIENT OVERLAP, AND GRADIENT NORMS

These raw-gradient diagnostics use the M-PG/100 student, whereas Figure 3(b,c) measures teacher-conditioned optimizer steps at M-PG/500. Batches use seeds 42 and 43, one response per domain per batch, a 2,048-token prompt cap, and a 256-token generation cap. All eight responses reach the cap. Four response positions are sampled without replacement per response. Each response’s four prefixes receive equal weight. For domain-specific inputs, each teacher scores the response from its own domain; for shared inputs, each teacher scores all four responses with equal domain weights. A third control evaluates every teacher on every input domain separately, keeping four prefixes per gradient. These conditions reuse the same responses, allowing paired comparisons within each batch.

Let $S_d(\rho)$ contain the largest-magnitude ρ fraction of teacher d ’s gradient or update coordinates, with stable coordinate-index tie breaking. Jaccard overlap $J(d, e) = |S_d(\rho) \cap S_e(\rho)| / |S_d(\rho) \cup S_e(\rho)|$ measures shared locations; cosine also captures signs and magnitudes (Yu et al., 2026). Table 8 reports PG gradient norms for domain-specific inputs; the first batch’s instruction-following gradient has norm 6.13×10^{-9} . Cosine measures directional opposition, while the gradient norm determines its magnitude. Independent uniform sets of fraction ρ have expected-intersection to expected-union ratio $\rho / (2 - \rho)$, approximately 0.00503 at $\rho = 1\%$. This reference uses uniform coordinates; layer and weight-magnitude structure remain in the measured gradients.

The largest-magnitude 1% of raw-gradient coordinates give mean Jaccard overlaps of 0.0951 for PG and 0.0999 for I64 on domain-specific inputs, averaging the two batches and six teacher pairs. Their ranges are 0.0744–0.1077 and 0.0737–0.1178. When every teacher scores all four domain responses, the respective means are 0.3808 and 0.4577. The latter comparison changes both input composition and prefix count; the teacher-by-input-domain comparison below fixes the prefix count.

Holding the teacher fixed while changing the input domain gives mean gradient cosines of 0.0059 for PG and 0.0072 for I64. Holding the input domain fixed while changing teachers produces both aligned and opposed gradients. For domain-specific inputs the mean cosines are -0.0020 and 0.0081 , respectively, alongside partial coordinate overlap. Figure 8 shows how input choice affects the teacher-overlap measurements in Section 4.

D VOCABULARY SUPERVISION AND COVERAGE

Single-step comparison using four response batches. The M-PG/100 row of Figure 4(a) uses four cached response batches with seeds 2042–2045. Each batch contains one response per domain. Generation uses the training prompt format, a 2,048-token prompt cap, and a 4,096-token response cap. The 16 response lengths range from 81 to 4,096 tokens; two responses reach the cap. Six positions per response give 96 prefixes. PG actions are independently sampled from the student at these fixed prefixes. All losses use the same prefixes, equal domain weights, and equal response

Table 7: **Vocabulary supervision produces similar single-step update concentration.** M-PG/100 means over four cached batches. NZ: nonzero coordinates (%); top-1%: squared step norm in the largest 1% of coordinates (%).

	Loss	FP32 NZ	BF16 NZ	FP32 top-1%	BF16 top-1%
PG		87.982	0.149	8.845	100.000
I64		87.989	0.150	8.827	100.000
T64		87.989	0.150	8.828	100.000
Full		87.999	0.150	8.828	100.000

Table 8: **The teacher top-64 set captures nearly all student probability mass in these controls.** Means over sampled prefixes; gradient norms use domain-specific inputs. Entropy is in nats, and 100% may reflect rounding.

Bank	Domain	Mass in T (%)	Mass in I (%)	Entropy	$\ g_{PG}\ _2$
1	Math	100.0000	100.0000	0.3243	132.6
1	Code	99.9968	99.9965	0.7701	163.8
1	IF	100.0000	100.0000	1.85×10^{-5}	6.13×10^{-9}
1	Science	99.9184	99.9162	0.7405	54.82
2	Math	99.9991	99.9991	0.3221	2.605
2	Code	100.0000	100.0000	0.2962	3.167
2	IF	100.0000	100.0000	0.5004	55.52
2	Science	99.9993	99.9992	0.2442	0.8505

weights within each domain. Gradients are globally clipped at norm 1 before applying the same saved Adam state. Table 7 adds the FP32 and energy-concentration measurements for this state.

The loss and Adam implementations match the two-batch controls in Appendix C. Appendix G.3 reports the actual I64 intersection coverage at all four student checkpoints.

Implementation details. For I64 (Equation 3), candidate sets and teacher scores are fixed at roll-out. Each student forward recomputes current probabilities at the saved candidate positions and refreshes the detached coefficients. PG also refreshes its detached log-ratio coefficients and uses no advantage or policy-ratio clipping.

Teacher-top-64 reference objective. The teacher-top-64 diagnostic uses the fixed teacher set T and full-vocabulary probabilities:

$$\ell_{T64} = \sum_{v \in T} \left[p_{\theta}(v) \log \frac{p_{\theta}(v)}{q_d(v)} - p_{\theta}(v) + q_d(v) \right]. \quad (9)$$

Its gradient is $\sum_{v \in T} p_{\theta}(v) \log(p_{\theta}(v)/q_d(v)) \nabla_{\theta} \log p_{\theta}(v)$. I64 uses the detached weights in Equation 3, normalized over the original student set S . T64 and full-vocabulary KL are evaluated in single-step comparisons. End-to-end training compares PG, I16, and I64; the I16 measurements cover task performance and online coverage.

Differentiating Equation 2 gives

$$\nabla_{\theta} \ell_{\text{full}} = \sum_{v \in \mathcal{V}} p_{\theta}(v) \left(\log \frac{p_{\theta}(v)}{q_d(v)} + 1 \right) \nabla_{\theta} \log p_{\theta}(v) = \mathbb{E}_{a \sim p_{\theta}} \left[\log \frac{p_{\theta}(a)}{q_d(a)} \nabla_{\theta} \log p_{\theta}(a) \right] \quad (10)$$

using $\sum_v \nabla_{\theta} p_{\theta}(v) = 0$. This identity conditions on fixed prefixes and weights; the training loss additionally depends on the sampled response and the token weights induced by its final length and the normalization rule. Table 8 reports candidate-set coverage for the two-batch gradient-overlap and Adam-state controls.

E SUPPORTING MEASUREMENTS FOR THE MAIN FIGURES

Teacher-update controls. Figure 3(a) compares cumulative parameter changes at update 500 under DR, with additional I64-DT/GT rows; paired points distinguish FP32 and BF16 for each statistic.




Figure 9: **Mean performance gains can coexist with a regression in an individual domain.** (a) Four-domain mean $\pm$ std; S/T mark the initial student/teacher. (b) Worst-domain change from initialization; zero means no regression. Lines connect reported point estimates.

Panels (b,c) start from the same M-PG/500 student parameters and Adam moments. With domain-specific inputs, each teacher scores responses from its own domain; with shared inputs, all teachers score the same prefixes. Each of four cached response batches contains six teacher pairs. Small points summarize pairs within a batch, and large points average batches. Equation 6 subtracts the Adam step with the current gradient set to zero before computing FP32 update cosines. The resulting step difference remains conditional on the saved moments through Adam’s nonlinear update. SGD and reset- m steps are individually rescaled to the norm of the Adam step with trained moments; reset- m retains the second moment and step counter. Teacher pairs reuse teachers and responses, so pairwise observations are dependent. These calculations use non-fused PyTorch updates on copies of saved checkpoints; the online training optimizer uses a fused implementation, allowing numerical differences between them. Mathematics-only and joint training also differ in mathematics prompt exposure (Appendix A).

Vocabulary-supervision controls. Figure 4(a,b) uses four cached response batches at each of four student checkpoints: M-PG and M-I64-DR at updates 100 and 500. Losses use identical prefixes and Adam moments within each checkpoint, while responses can differ between checkpoints. Panel (a) reports batch means. In panel (b), points are individual batches, lines connect PG batch means, and shading spans the I64 measurements. The direction measurements cover PG, I64, and full-vocabulary supervision; I16 appears in the end-to-end training results. Panel (c) compares S-I64, M-I64, and M-I16 with PG at the same task, teacher setting, and update. Table 7 adds FP32 and energy-concentration measurements for the M-PG/100 comparison in Appendix D.

E.1 ADDITIONAL OPTIMIZER AND SAMPLING CONTROLS

Figure 2 uses the six cached response batches and training clipping records described in Appendix B.1. In panel (a), boxes mark the interquartile range and whiskers the 5th–95th percentiles; labels give the fraction of training updates clipped. Reset- m and reset- (m, v) retain the optimizer step counter. Panel (c) averages four domain probe losses within each batch; error bars show sample standard deviations over the three batches at each checkpoint.

Figure 6(a) uses three M-PG/100 response batches, with a 4,096-token response cap, containing 48 responses and 288 prefixes. Twelve fixed subsets of prefixes each receive two independent draws of 1, 16, or 64 PG actions per prefix. I64 is evaluated once per subset. Increasing the sample count raises the mean PG gradient cosine to full-vocabulary KL from 0.5374 to 0.8652 and 0.9667; I64 reaches 0.999988. These repeats measure sampling variation at a fixed checkpoint. Panels (b,c) use the separate two-batch experiment with a 256-token response cap (Appendix C); its zero-gradient control is evaluated once. Reset- m retains the second moment and step counter, whereas fresh Adam resets both moments and the counter. The sampling and moment-reset comparisons each use their own response batches.

Cumulative parameter changes. Figure 10 shows the trajectories underlying the update-500 summary in Figure 3(a).




Figure 10: **Cumulative BF16 changes remain much sparser than FP32 changes.** PG/I64 trajectories under DR in single-teacher (S) and joint (M) training, with joint I64-DT/GT added: (a,b) nonzero coordinate fractions; (c,d) coordinate fractions containing 90% of squared displacement from initialization. Lines connect measured updates 1, 50, 100, 250, and 500; dashed lines denote S. The curves describe the recorded trajectories; uncertainty across training seeds remains unquantified.

F PROMPT EXPOSURE AND CANDIDATE-SET WEIGHTS

When each domain loss is averaged separately, its prompt quota and loss weight control different quantities. Increasing the quota at fixed λ_d changes the number of sampled responses and the variance of the gradient estimate, without directly increasing the domain’s expected contribution to the total gradient. Mathematics-only runs receive four times as many mathematics prompts per update as joint runs. Their score differences reflect changes in both teacher count and mathematics exposure.

Candidate-set coverage also scales the distillation gradient. For the student candidate set S and its intersection I with the teacher set in Equation 3, the retained student weights sum to $r(h) = p_\theta(I | h)/p_\theta(S | h)$. At a fixed prefix with a nonempty intersection and detached coefficients, the I64 gradient equals $r(h)$ times the gradient obtained by renormalizing the same weights over I . Coverage therefore changes gradient scale across prefixes even under equal response weighting. Measuring the retained mass and empty-intersection rate by domain can reveal unequal coverage under a fixed candidate size. Enlarging the candidate set or renormalizing its intersection changes the objective and requires a separate performance comparison.

Table 9: **Method rankings depend on the task and test set.** Additional test-set scores (%) for the same four domains; mean weights domains equally.

Configuration	Update	Math	Code	IF	Science	Mean
Initial student	0	42.97	23.44	25.78	62.21	38.60
PG-DR / Adam	500	46.48	26.17	37.50	64.06	43.55
PG-DT / Adam	250	44.53	26.56	36.72	64.75	43.14
PG-DT / Adam	500	44.92	27.34	38.28	63.48	43.51
PG-GT / Adam	500	44.92	28.91	39.06	65.14	44.51
I16-DR / Adam	500	44.92	28.52	40.23	64.55	44.56
I64-DR / Adam	500	46.88	26.17	39.06	65.72	44.46
I64-DT / Adam	500	43.36	27.73	39.06	64.55	43.68
I64-GT / Adam	500	45.31	27.34	41.02	66.21	44.97
PG-DR / SGD	500	46.88	27.73	37.11	62.21	43.48
PG-DT / SGD	500	46.48	28.52	37.11	63.48	43.90
PG-GT / SGD	500	44.53	26.17	36.72	64.45	42.97

G ADDITIONAL TEST RESULTS AND ROBUSTNESS CHECKS

G.1 ADDITIONAL TEST SETS FOR THE SAME FOUR DOMAINS

The ten joint-training configurations use the development benchmarks and checkpoint grid in Appendix A. The following analyses evaluate selected checkpoints on separate test questions, measure candidate-set coverage, and check sensitivity to known training–evaluation overlap.

Table 9 uses additional test sets fixed before inference, with 256 questions per domain: 1,024 questions and 1,792 responses per model. Mathematics and code come from the Skywork-OR1-RL-Data mathematics and TACO-based code pools; instruction following uses Nemotron-RL-instruction_following; science uses the Nemotron-RL-knowledge-mcqa validation split restricted by a fixed natural-science keyword filter. Questions are the first eligible unique candidates in the fixed source order, selected without model scores. Normalized exact-match and lexical near-duplicate checks remove questions matching the original training bundle or known evaluation pools. The audit covers exact and lexical overlap with these pools; semantic overlap and pretraining exposure remain outside its scope.

Scoring uses mathematics answer verification, code unit tests, instruction-following constraints, and science multiple-choice accuracy. Mathematics, code, and instruction following use greedy decoding; science uses four responses per question at temperature 1. The context budget is 32,768 tokens, with thinking disabled. The data distributions and code/instruction-following scorers differ from the development benchmarks, so their scores are reported separately. The four-domain mean weights domains equally. Question-paired bootstrap intervals quantify finite-test uncertainty conditional on the trained models.

G.2 CHECKPOINT SELECTION BEFORE TEST EVALUATION

We apply the three selection rules in Section 6 to the development grid in Appendix A. The grid contains updates 50/100/250/500 for seven Adam trajectories and 100/250/500 for three SGD trajectories. The retention tolerance is two percentage points per domain relative to the initial student, fixed before the additional test results. Both development rules keep the initial student unless the selected equal-domain mean improves by at least one percentage point. Checkpoint choices are fixed before test inference and remain unchanged after the test and overlap-sensitivity analyses. Both development-based rules choose update 250 for PG-DT/Adam and update 500 for the other nine trajectories. The separate PG-DT/Adam update-250 test evaluation allows direct comparison with its update-500 endpoint.

Table 10 summarizes the resulting checkpoint choices and test-score differences.

Table 11 gives all 37 trained-checkpoint development evaluations, including the ten endpoints already reported in Table 2. These intermediate measurements evaluate checkpoints along the same ten training trajectories.

Table 10: **The retention constraint leaves every checkpoint choice unchanged.** Test Δ is the selected-checkpoint mean minus the update-500 mean (percentage points).

Trajectory	Endpoint	Max dev. mean	Retention	Test Δ (pp)
Other nine trajectories	500	500	500	0.000
PG-DT / Adam	500	250	250	-0.366

The retention-constrained and maximum-mean rules select the same checkpoint in every trajectory. On the additional test sets, the update-250 PG-DT/Adam checkpoint scores 0.366 percentage points below its update-500 endpoint in four-domain mean. The endpoint therefore has the higher separate-test mean in this comparison. All selected models satisfy the retention and useful-gain conditions at the level of test point estimates; joint statistical uncertainty across domains remains unquantified. PG-DR/Adam has two intermediate development checkpoints outside the retention tolerance: GPQA falls from 32.20% initially to 28.03% at update 50 and 30.18% at update 250, before reaching 35.97% at update 500. The learning curve captures these temporary regressions and their subsequent recovery.

G.3 PROBABILITY MASS RETAINED BY THE CANDIDATE INTERSECTION

For each of four student checkpoints (M-PG and M-I64-DR at updates 100 and 500), four cached response batches provide one response per domain and six sampled positions per response: 96 prefixes per checkpoint. Each response is scored by its domain teacher. We measure $p_\theta(I | h)$ for the actual intersection $I = S \cap T$ and the retained fraction $r = p_\theta(I | h) / p_\theta(S | h)$.

The minimum student probability mass in the actual student–teacher top-64 intersection is 96.55% over the 384 sampled prefixes, and no intersection is empty (Table 12). Thus the single-step I64/full-vocabulary comparisons use prefixes where truncation discards little student probability mass. The four batches measure response-sampling variation at each checkpoint, with potentially different responses between checkpoints; the training trajectory is held fixed.

G.4 SENSITIVITY TO KNOWN MATHEMATICS OVERLAP

The fixed Qwen training pool contains five records matching four known MATH-500 problems, including variants in wording or answer format. The affected MATH test identifiers are counting-and-probability problem 430 and intermediate-algebra problems 1849, 2146, and 582. The audit identifies training-pool membership; run-specific exposure and its effect on scores remain unmeasured. We remove these four questions from the existing per-question evaluation results, retaining the original 500-question results as the main protocol.

Removing the four questions changes the four-domain means by approximately 0.09–0.16 points and can reorder near-tied configurations. PG-DT/Adam and PG-GT/Adam both round to 38.60% before removal, but reach 38.70% and 38.75%, respectively, afterward. PG-DR/SGD retains the highest mean among the reported checkpoints. MATH-496 denotes this retained subset of MATH-500. Checkpoint choices and trained weights remain fixed, so the analysis measures score sensitivity to the four known overlaps.

G.5 TEACHER-UPDATE OVERLAP AND DISTRIBUTIONAL SIMILARITY

Figure 11 supplements the zero-gradient baseline subtraction at M-PG/500 using the controls in Appendix E. It compares PG and I64 under Adam with trained moments, reset- m , and norm-matched SGD. Shared inputs hold student prefixes fixed between teachers; domain-specific inputs change both the teacher and the input domain. Each of four batches contains six teacher pairs; these observations are dependent because they reuse teachers and responses. These measurements characterize local gradient and update geometry; task performance is evaluated separately.

Table 11: **Complete joint-training checkpoint scores.** All 37 evaluated checkpoints from ten Qwen3-1.7B joint-training configurations. Scores (%) use the original MATH-500, LiveCodeBench, IFBench strict, and GPQA avg@4 protocol; Mean is the unweighted average of the four domains. † marks the ten update-500 rows also reported in Table 2.

Configuration	Update	MATH	Code	IF	GPQA	Mean
PG-DR / Adam	50	69.60	16.41	18.67	28.03	33.18
	100	72.00	17.19	19.33	31.19	34.93
	250	73.40	16.41	26.00	30.18	36.50
	500†	73.20	18.27	27.00	35.97	38.61
PG-DT / Adam	50	71.60	17.19	19.00	34.34	35.53
	100	74.60	19.53	18.00	34.97	36.78
	250	77.00	19.53	25.00	33.59	38.78
	500†	73.20	19.24	27.00	34.96	38.60
PG-GT / Adam	50	71.20	18.75	19.00	34.72	35.92
	100	74.80	21.88	17.00	31.94	36.41
	250	74.80	19.53	25.00	33.84	38.29
	500†	76.20	17.92	27.33	32.94	38.60
I16-DR / Adam	50	69.80	17.19	17.33	31.94	34.07
	100	72.20	17.97	19.00	32.95	35.53
	250	75.60	18.75	25.67	33.59	38.40
	500†	75.40	17.75	27.00	35.46	38.90
I64-DR / Adam	50	69.60	17.19	19.00	31.69	34.37
	100	71.60	17.97	19.00	32.70	35.32
	250	74.80	17.97	26.00	33.96	38.18
	500†	75.80	18.30	27.33	34.71	39.04
I64-DT / Adam	50	72.20	18.75	17.00	32.32	35.07
	100	74.80	19.53	19.33	33.33	36.75
	250	75.40	21.09	25.33	32.70	38.63
	500†	74.00	18.86	28.67	35.22	39.19
I64-GT / Adam	50	74.20	17.97	17.00	32.95	35.53
	100	74.40	17.97	21.00	32.07	36.36
	250	74.40	18.75	25.00	33.33	37.87
	500†	74.20	19.80	28.67	33.41	39.02
PG-DR / SGD	100	73.00	18.75	25.67	35.35	38.19
	250	75.80	18.75	25.67	31.82	38.01
	500†	76.40	19.53	28.33	35.59	39.96
PG-DT / SGD	100	74.80	18.75	23.33	34.34	37.81
	250	74.20	19.53	25.33	35.10	38.54
	500†	75.20	20.31	28.33	33.90	39.44
PG-GT / SGD	100	74.60	17.97	20.67	35.73	37.24
	250	73.00	20.31	23.67	33.33	37.58
	500†	75.00	19.26	25.33	34.71	38.58

Table 12: **The minimum observed I64 intersection mass is 96.55%.** Statistics over sampled prefixes; $p(I)$ is in percent and $r = p_\theta(I | h)/p_\theta(S | h)$.

Checkpoint	$p(I)$: mean	$p(I)$: min	$ I $: mean	$ I $: min	r : mean	r : min
M-PG/100	99.9141	96.5482	58.96	43	0.9999413	0.9977714
M-PG/500	99.9952	99.7825	61.59	50	0.9999975	0.9998511
M-I64-DR/100	99.9564	97.5664	60.08	36	0.9997331	0.9757993
M-I64-DR/500	99.9826	99.3722	61.75	51	0.9999912	0.9997005

Table 13: **Removing four overlapping mathematics questions can reorder near-tied configurations.** Mean (496) substitutes MATH-496 for MATH-500 and keeps the other domain scores fixed.

Configuration	Update	MATH-500	MATH-496	Mean (500)	Mean (496)
Initial student	0	68.80	69.15	33.94	34.02
PG-DR / Adam	500	73.20	73.79	38.61	38.76
PG-DT / Adam	250	77.00	77.42	38.78	38.89
PG-DT / Adam	500	73.20	73.59	38.60	38.70
PG-GT / Adam	500	76.20	76.81	38.60	38.75
I16-DR / Adam	500	75.40	76.01	38.90	39.06
I64-DR / Adam	500	75.80	76.41	39.04	39.19
I64-DT / Adam	500	74.00	74.40	39.19	39.29
I64-GT / Adam	500	74.20	74.60	39.02	39.12
PG-DR / SGD	500	76.40	77.02	39.96	40.12
PG-DT / SGD	500	75.20	75.81	39.44	39.59
PG-GT / SGD	500	75.00	75.60	38.58	38.73




Figure 11: **Coordinate overlap and update direction measure different properties.** (a) Top-1% FP32 update-set Jaccard versus cosine; small/large points show batch/overall means. (b) Teacher-distribution JS divergence versus 1-Jaccard on shared inputs; each point is a teacher pair within a batch.

G.6 ONLINE COVERAGE AND PER-DOMAIN NORMALIZATION EFFECTS

Online I16 coverage. For I16-DR, each of 500 updates has 16 logged slices of four responses. Each update averages the masked, response-averaged summaries with equal weight per slice. Candidate sets are fixed at rollout; student probabilities are measured during training. Across 8,000 slices, mean $p_\theta(I)$, r , and $|I|$ are 99.616%, 99.879%, and 15.048. The mean empty fraction is 0.00270%, nonzero in 80 slices across 62 updates. The lowest slice-mean mass is 86.67% and largest slice-mean empty fraction 1.85%; these extremes refer to slice means.

Per-domain probe losses. We disaggregate Figure 2(c), using the same six batches and twelve branches per batch (Appendix B.1). Each batch has one held-out response per domain. At update 250, saved-state Adam under DR/DT/GT increases IF KL by $0.362/0.293/0.272 \times 10^{-3}$, while the four-domain changes are $0.041/-0.002/-0.027 \times 10^{-3}$. Thus the mean reduction under DT/GT coexists with an instruction-following increase.




Figure 12: **High average coverage and low mean probe loss can conceal exceptions.** (a–c) I16 per-update means of 16 logged slices, without temporal smoothing; the mass axis spans 98.5–100%. (d,e) Full-vocabulary KL after minus before BF16 writeback, in units of 10^{-3} . Cells show three-batch means; dots show batches 1042–1044 from left to right. The shared symmetric-log color scale is linear within ± 0.1 . SGD is matched to each rule’s Adam step norm.

H SMOLLM3 DIAGNOSTICS AND CAPABILITY EVALUATION

H.1 CHECKPOINTS, RESPONSES, AND WEIGHTING

We use the SmolLM3-3B MixSFT student, three domain-specific RL teachers, and exported 500-update students for I64-DR/DT/GT and PG-DR/SGD. The model and data revisions, export hashes, and diagnostic source snapshots are recorded with the measurements. The reported I64 training configurations use Adam with learning rate 2.5×10^{-7} ; PG uses SGD with learning rate 0.0025. The capability comparison uses a new shared evaluation protocol. Recovery of the three original I64 training source revisions and prompt schedules is incomplete, which limits attribution of endpoint differences to individual training choices.

Diagnostic prompts come from the frozen Open-MOPD data revision 9e897efe, with unique raw messages shuffled within mathematics, code, and instruction following. Coverage, normalization, and probe prompt splits are disjoint within this diagnostic sample. We preserve dataset messages and disable thinking in the chat template. Generation uses temperature 1, top- $p = 1$, a 2,048-token prompt cap, and a 4,096-token response cap. Each natural coverage bank contains 384 responses, 128 per domain, with up to eight uniformly spaced positions per response: 3,053 prefixes for MixSFT and 3,058 for DR500. Additional students score the same 3,058 DR500 prefixes without regenerating responses.

Coverage statistics assign equal weight to domains, responses within domains, and audited positions within a response. Mean-mass intervals resample responses independently within each domain 2,000 times using seed 1042/1043/1044. The first two rows of Table 14 use different natural banks; the remaining students use the DR500 bank. The minimum I64 mass on DR500’s natural bank is 1.7454×10^{-8} ; DT500 and GT500 each have one empty intersection in the common-bank audit.

Table 14: **High average I64 coverage coexists with a lower tail.** SmolLM3, 384 responses per bank. Mass and frequency columns are percentages; intervals resample responses within domains. The last column counts empty intersections among audited prefixes.

Student	Response bank	Mean mass [95% CI]	< 90%	1st pct.	Empty
MixSFT	MixSFT	97.42 [96.93, 97.88]	8.42	62.06	0
DR500	DR500	98.56 [98.16, 98.96]	3.65	64.75	0
GT500	DR500	98.54 [98.13, 98.94]	3.61	64.79	1
DT500	DR500	98.58 [98.18, 98.97]	3.52	65.41	1
PGSGD500	DR500	98.65 [98.27, 99.03]	3.52	66.49	0

H.2 FROZEN SUBSETS AND GRADIENT COMPARISONS

Within each domain we select representative prefixes, prefixes above its 75th entropy percentile, and prefixes with I64 mass below 0.99. Selection precedes the objective-gradient comparison, allows overlap between strata, and uses up to 32 distinct prompts per domain. The low-coverage code stratum has 21 eligible prompts, yielding 85 distinct prefixes overall. Each stratum is divided into three batches. Small candidate budgets $k = 4, 8$ are selected by a coverage-only rule before inspecting objective gradients.

The BF16 coverage audit sorts a top-128 list and takes its first k entries, whereas gradient diagnostics recompute top- k directly in FP32 at the exported BF16 parameter values. Student forward/backward passes use FP32; cached teacher targets use BF16 forward passes followed by FP32 log-softmax. These precision and top- k choices can change probabilities and threshold membership. One of the 85 frozen low-coverage prefixes crosses 0.99 during the FP32 diagnostic and remains in its original stratum. The FP32 diagnostic mean I64 masses are 99.049%, 95.684%, and 93.121% for the representative, high-entropy, and low-coverage groups. These domain-weighted summaries characterize the selected prefixes; natural-generation prevalence is measured by the coverage audit.

Full KL, PG, and intersection objectives use the definitions in Appendix D. Empty intersections produce zero loss. We retain zero- and small-norm cases and report vector differences in addition to cosine. The plotted relative error is $\|g_{I_k} - g_{\text{full}}\|/\|g_{\text{full}}\|$, which measures the vector difference

relative to the full-vocabulary gradient norm. Batch means and thin batch trajectories in Figure 5 make the diagnostic replication unit explicit.

H.3 OPTIMIZER STATE, SAMPLING, AND NUMERICAL PRECISION

Normalization uses three separate batches, each with eight responses per domain, and the true response lengths for GT/DT weights. Raw gradients, fresh Adam, and SGD start from independent copies of the same exported parameters. Fresh Adam starts with both moments zero and performs its first step; SGD uses learning rate 0.0025, zero momentum, and norm-1 clipping. Fresh Adam is distinct from the Qwen reset- m control, which retains the trained second moment and step counter. These controls use their specified learning rates; a native Adam step for norm matching is unavailable.

For all three rule pairs in all three batches, fresh Adam lowers cosine relative to the raw gradients. For example, DT-DR raw cosines are 0.8968/0.8655/0.5933, versus 0.7013/0.6489/0.4126 after fresh Adam.

Conditional PG variance uses only three fixed prefixes per group, one per domain, with eight independent action-sampling repetitions. Sample budgets $m = 1, 16, 64$ use nested actions within a repetition. We normalize the total coordinate sample variance by $\|g_{\text{full}}\|^2$ and separately report the error of the eight-repetition mean. The representative-group variance is 1.426/0.570/0.123; the low-coverage-group variance is 4.965/2.566/0.434.

The precision controls begin exactly at parameter values represented by BF16, cast to FP32 for computation. Native DR500 FP32 master weights, Adam moments and step, and the initialization master are unavailable for these diagnostics.

H.4 CUMULATIVE BF16 CHANGES BY PARAMETER GROUP

The exported MixSFT-to-DR500 comparison contains 3,075,098,624 unique parameter coordinates after accounting for tied embeddings. The 36 transformer blocks each contain 78,123,008 coordinates, the token embedding 262,668,288, and the final normalization 2,048. We compare the two BF16 exports directly. Overall, 10.4695% of coordinates change; the largest 1% carry 36.8477% of squared displacement, and 5.8496% of coordinates carry 90%. Figure 14 retains all parameter groups. Per-group changed fractions use each group’s own coordinate count; displacement shares use total model squared displacement. These statistics describe displacement between BF16 exports; native FP32 history is unavailable.




Figure 13: **Local controls separate optimizer, sampling, and precision effects.** (a,b) Three paired SmolLM3 batches (thin) and their means (bold). In (b), x is the gradient or step, normalized by the first named rule. (c,d) Three fixed prefixes per group, eight action repeats; (d) error of the vector mean. (e,f) Each rule/batch pairs FP32 changes with BF16 writeback. Fresh Adam starts at the exported BF16 parameter values with both moments zero.




Figure 14: **Cumulative BF16 changes differ across all 38 SmolLM3 parameter groups.** MixSFT to DR500: blocks 0–35, embedding (E), final norm (N). (a) Within-group changed fraction; dashed: global parameter-weighted fraction. (b) Share of total squared displacement. Tied weights count once; native FP32 changes are unavailable.

H.5 ENDPOINT CAPABILITY EVALUATION

Table 15 evaluates the exported students with shared prompts, scorers, and decoding settings within each benchmark. We use BF16, disable thinking, and allocate a 32,768-token context with the response budget equal to context length minus prompt length. The SmolLM evaluation uses the separately recorded 128-question LiveCodeBench subset; the main Qwen study uses 1,055 questions. MATH-500, this LiveCodeBench subset, and IFBench use greedy decoding; GPQA uses temperature 1 and top- $p = 1$; AIME uses temperature 0.6 and top- $p = 0.95$. Budget-capped outputs remain in scoring. Score and paired-difference intervals use 100,000 question-bootstrap replicates (seed 1042/1043/1044), retaining all responses to each sampled question.

Table 15: **SmolLM3 capability under a shared evaluation protocol.** Scores (%): AIME pass@8, GPQA average over four responses, otherwise greedy accuracy. $Q \times R$: questions $\times$ responses; Resp.: average response accuracy; Tokens: mean generated length; Cap: budget-capped responses (%). Intervals resample questions.

Benchmark	$Q \times R$	Score [95% CI]	Resp.	Tokens	Cap (%)
<i>MixSFT initialization</i>					
MATH-500 [†]	500×1	50.60 [46.20, 55.00]	50.60	1501	3.00
AIME 2024	30×8	10.00 [0.00, 23.33]	2.92	3840	5.83
AIME 2025	30×8	13.33 [3.33, 26.67]	1.67	2735	3.75
<i>I64-DT/Adam, update 500</i>					
AIME 2024	30×8	26.67 [13.33, 43.33]	9.58	5019	7.50
AIME 2025	30×8	13.33 [3.33, 26.67]	3.75	1227	0.00
<i>PG-DR/SGD, update 500</i>					
MATH-500	500×1	53.80 [49.40, 58.20]	53.80	1520	2.80
LiveCodeBench v6	1055×1	17.97 [11.72, 25.00]	17.97	1920	5.47
IFBench (strict)	300×1	20.33 [16.00, 25.00]	20.33	2066	5.33
GPQA Diamond	198×4	23.99 [20.71, 27.40]	23.99	543	0.00
AIME 2024	30×8	16.67 [3.33, 30.00]	4.17	2215	1.67
AIME 2025	30×8	13.33 [3.33, 26.67]	2.08	4957	9.17

[†] Additional complete MixSFT run under the same protocol (500 scored responses).

MATH-500 greedy accuracy is 50.60% for MixSFT and 53.80% for update-500 PG-DR/SGD, a change of +3.20 points (95% paired interval $[-0.20, 6.60]$), gaining 45 questions and losing 29. Update-500 I64-DT/Adam changes AIME 2024 pass@8 by +16.67 points ($[0.00, 33.33]$) relative to MixSFT.

Equal endpoint scores can hide different solved questions. AIME pass@8 is the fraction of questions with at least one correct answer among eight samples. MixSFT, update-500 I64-DT/Adam, and update-500 PG-DR/SGD solve 3/30, 8/30, and 5/30 questions on AIME 2024, respectively. On AIME 2025, each solves 4/30. Their average response accuracies are 1.67/3.75/2.08%, and only one solved question is common to all three. MixSFT and PG-DR/SGD also share only one of their four solved questions (Figure 15). These endpoint comparisons describe aggregate and question-level outcomes across students whose training choices vary together. Attribution to individual choices remains limited by the recovered training records.




Figure 15: **Equal aggregate scores can conceal different solved questions and sampling outcomes.** (a) Differences from MixSFT with paired question-bootstrap 95% intervals: accuracy for MATH-500, pass@8 for AIME. (b) Correct responses out of eight for every AIME 2025 question solved by any model; the other 22 questions have zero successes. (c) Subset estimates from the same eight responses: $\widehat{\text{pass@k}} = \frac{1}{30} \sum_q [1 - \binom{8-c_q}{k} / \binom{8}{k}]$, where c_q counts correct responses.